

Users, Files, and the Manual

Prof. Seokin Hong

Kyungpook National University

Fall 2018

Objectives

■ Ideas and Skills

- The role and use of on-line documentation
- The Unix file interface: open, read, write, lseek, close
- Reading, creating, and writing files
- File descriptors
- Buffering: user level and kernel level
- Kernel mode, user mode, and the cost of system calls
- How Unix represents time, how to format Unix time
- Using the utmp file to find list of current users
- Detecting and reporting errors in system calls

■ System Calls and Functions

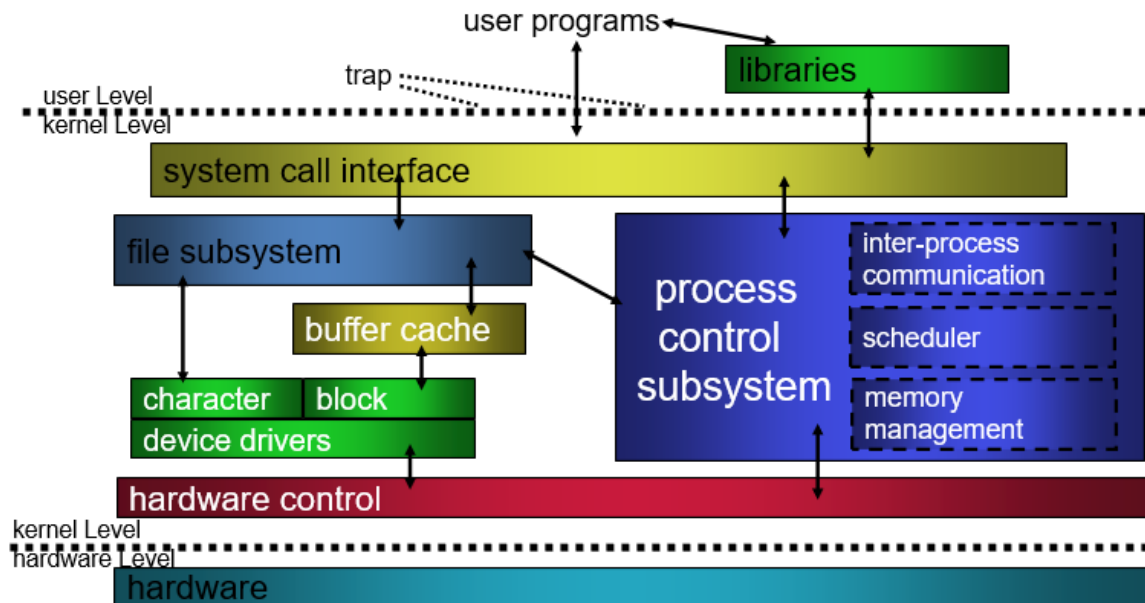
- open, read, write, create, lseek, close
- perror

■ UNIX/LINUX Commands

- man, who, cp, login

System Call

- System calls are an operating system's API
 - The set of functions that the operating system exposes to processes
- If you want to the OS to do something, you tell it via a system call



What are System Calls used for?

- Anything to do with
 - Accessing devices
 - Accessing files
 - Requesting memory
 - Setting/Changing access permissions
 - Communicating with other processes
 - Stopping/starting processes
 - Setting a timer

- You need a system call to
 - Open a file
 - Get data from the network
 - Kill a process

Three step to learn Unix/Linux system programming

- Looking at “real” programs : What does that do?
- Looking at the system calls : How does that work?
- Writing our own version : Can I try to do it?

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Asking About **who**

- What does **who** do?
- How does **who** work?
- Can I write **who**?




FIGURE 2.1

Users, files, processes, devices, and kernel.

Commands are Programs

- Almost all Unix commands are simply programs written by a variety of people, usually in C.
- When you type **ls**, you are asking your shell to run the program named **ls**.
- Adding new commands to Unix is easy;
You write a new program and have the executable file stored in one of the standard system directories:
/bin, /usr/bin, /usr/local/bin.

What does **who** do?

Q1: What Does **who** Do?

- The **who** command displays information about users and processes on the local system

```
$ who
```

hecker1	ttyp1	Jul 21 19:51	(tide75.surfcity.com)
nlopez	ttyp2	Jul 21 18:11	(roam163-141.student.ivy.edu)
dgsulliv	ttyp3	Jul 21 14:18	(h004005a8bd64.ne.mediaone.net)
ackerman	ttyp4	Jul 15 22:40	(asd1-254.fas.state.edu)
wwchen	ttyp5	Jul 21 19:57	(circle.square.edu)
barbier	ttyp6	Jul 8 13:08	(labpc18.elsie.special.edu)
ramakris	ttyp7	Jul 13 08:51	(roam157-97.student.ivy.edu)
czhu	ttyp8	Jul 21 12:47	(spa.sailboat.edu)
bpsteven	ttyp9	Jul 21 18:26	(207.178.203.99)
molay	ttypa	Jul 21 20:00	(xyz73-200.harvard.edu)

\$

username

terminal name

login time

login host

one log-in session

What does **who** do?

Reading the Manual

- Every Unix system comes with manual for all commands
- The manual is on the disk and the command to read a page from the manual is **man**

```
$ man who
who(1)

NAME

    who - Identifies users currently logged in

...
SYNOPSIS

    who [-a] | [-AbdhHlmMpqrstTu] [file]

    who am i

    who am I

    whoami
```

Q2: How Does **who** Do It?

- To learn more about Unix from Unix
 - Read the manual
 - Search the manual
 - Read the .h files
 - Follow SEE ALSO links

Read the Manual

- `$ man who`

DESCRIPTION

The `who` utility can list the user's name, terminal line, login time, elapsed time since activity occurred on the line, and the process-ID of the command interpreter (shell) for each current UNIX system user. It examines the `/var/adm/utmp` file to obtain its information. If file is given, that file (which must be in `utmp(4)` format) is examined. Usually, file will be `/var/adm/wtmp`, which contains a history of all the logins since the file was last created.

Search the Manual

■ To find more about utmp

○ \$ man -k utmp

getutxid	getutxent (3c)	- access utmpx file entry
getutxline	getutxent (3c)	- access utmpx file entry
pututline	getutent (3c)	- access utmp file entry
pututxline	getutxent (3c)	- access utmpx file entry
setutent	getutent (3c)	- access utmp file entry
setutxent	getutxent (3c)	- access utmpx file entry
ttyslot	ttyslot (3c)	- find the slot in the utmp file of the current user
updwtmp	getutxent (3c)	- access utmpx file entry
updwtmpx	getutxent (3c)	- access utmpx file entry
<u>utmp</u>	<u>utmp (4)</u>	<u>- utmp and wtmp entry formats</u>
utmp2wtmp	acct (1m)	- overview of accounting and miscellaneous accounting commands
utmpd	utmpd (1m)	- utmp and utmpx monitoring daemon
utmpname	getutent (3c)	- access utmp file entry
utmpx	utmpx (4)	- utmpx and wtmpx entry formats
utmpxname	getutxent (3c)	- access utmpx file entry
<u>wtmp</u>	<u>utmp (4)</u>	<u>- utmp and wtmp entry formats</u>
wtmpx	utmpx (4)	- utmpx and wtmpx entry formats
\$		

○ utmp (4) means the manpage for utmp is in section 4 of the manual

What does **who** do?

```
$ man 4 utmp
```

```
utmp(4)
```

```
utmp(4)
```

NAME

utmp, wtmp - Login records

SYNOPSIS

```
#include <utmp.h>
```

DESCRIPTION

The utmp file records information about who is currently using the system.

The file is a sequence of utmp entries, as defined in struct utmp in the utmp.h file.

The utmp structure gives the name of the special file associated with the user's terminal, the user's login name, and the time of the login in the form of time(3). The ut_type field is the type of entry, which can specify several symbolic constant values. The symbolic constants are defined in the utmp.h file.

The wtmp file records all logins and logouts. A null user name indicates a logout on the associated terminal. A terminal referenced with a tilde (~) indicates that the system was rebooted at the indicated time. The adjacent pair of entries with terminal names referenced by a vertical bar (|) or a right brace (}) indicate the system-maintained time just before and just after a date command has changed the system's time frame.

The wtmp file is maintained by login(1) and init(8). Neither of these pro-grams creates the file, so, if it is removed, record keeping is turned off. See ac(8) for information on the file.

FILES

/usr/include/utmp.h

/var/adm/utmp

more (88%)

What does **who** do?

Read the .h files

- \$ more /usr/include/utmp.h

```
#define UTMP_FILE      "/var/adm/utmp"
#define WTMP_FILE      "/var/adm/wtmp"

#include <sys/types.h>  /* for pid_t, time_t */

/*
 * Structure of utmp and wtmp files.
 *
 * Assuming these numbers is unwise.
 */

#define ut_name ut_user          /* compatibility */
struct utmp {
    char    ut_user[32];         /* User login name */
    char    ut_id[14];           /* /etc/inittab id- IDENT_LEN in
                                * init */
    char    ut_line[32];         /* device name (console, lnxx) */
    short    ut_type;            /* type of entry */
    pid_t    ut_pid;             /* process id */
    struct exit_status {
        short    e_termination; /* Process termination status */
        short    e_exit;         /* Process exit status */
    } ut_exit;                  /* The exit status of a process
                                * marked as DEAD_PROCESS.
                                */
    time_t    ut_time;           /* time entry was made */
    char    ut_host[64];         /* host name same as
                                * MAXHOSTNAMELEN */
};
/* Definitions for ut_type */
```

utmp.h (60%)

What does **who** do?

We now know how who works

- who works by:

```
open utmp  
[→ read record ←]  
[→ display record ←]  
close utmp→
```




FIGURE 2.2

Data flow in the who command.

utmp


Can I Write **who**?

- Two tasks we need to program
 - Read structs from a file
 - Display the information stored in a struct

Q: How do I read structs from a file?

■ Let's Read the Manual!

```
$ man -k file | pipe grep read
_llseek (2)          - reposition read/write file offset
fileevent (n)        - Execute a script when a channel becomes readable
                      or writable
gftype (1)           - translate a generic font file for humans to read
lseek (2)            - reposition read/write file offset
macsave (1)          - Save Mac files read from standard input
read (2)           - read from a file descriptor
readprofile (1)      - a tool to read kernel profiling information
scr_dump, scr_restore, scr_init, scr_set (3) - read (write) a curses
screen from (to) a file
tee (1)              - read from standard input and write to standard
                      output and files
$
```



Can I Write [who](#)?

\$ man 2 read

READ(2)	System calls	READ(2)
---------	--------------	---------

NAME

read - read from a file descriptor

SYNOPSIS

#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count);

DESCRIPTION

read() attempts to read up to count bytes from file descriptor fd into the buffer starting at buf.

RELATED INFORMATION (called SEE ALSO in some versions)

Functions: fcntl(2), creat(2), dup(2), ioctl(2), getmsg(2), lockf(3), lseek(2), mtio(7), open(2), pipe(2), poll(2), socket(2), socketpair(2), termios(4), streamio(7), opendir(3) lockf(3)

Standards: standards(5)

ANS: we use open, read, and close

- Opening a file: **open()**
 - Opening a file is a **kernel service**;
 - **open()** **system call** is a request from your program to the kernel

```
fd = open(name, mode)
      ↑      ↙
char *  O_RDONLY
        O_WRONLY
        O_RDWR

returns -1 on error
OR an int on success
```




FIGURE 2.3

A file descriptor is a connection to a file.

More about open()

open		
PURPOSE	Creates a connection to a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = open(char *name, int how)	
ARGS	name	name of file
	how	O_RDONLY, O_WRONLY, or O_RDWR
RETURNS	-1	on error
	int	on success

ANS: we use **open**, **read**, and **close**

- Reading Data from a File: **read()**

```
fd = open(name, mode)
```

```
n = read(fd, array, numchars)
```

read		
PURPOSE	Transfer up to qty bytes from fd to buf	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t numread = read(int fd, void *buf, size_t qty)	
ARGS	fd	source of data
	buf	destination for data
	qty	number of bytes to transfer
RETURNS	-1	on error
	numread	on success

ANS: we use `open`, `read`, and `close`

■ Closing a File: `close()`

```
fd = open(name, mode)
```

```
n = read(fd, array, numchars)
```

`close(fd)`

<code>close</code>	
PURPOSE	Closes a file
INCLUDE	<code>#include <unistd.h></code>
USAGE	<code>int result = close(int fd)</code>
ARGS	<code>fd</code> file descriptor
RETURNS	<code>-1</code> on error <code>0</code> on success

Writing who1.c

```
/* who1.c - a first version of the who program
 *
 *      open, read UTMP file, and show results
 */
#include      <stdio.h>
#include      <utmp.h>
#include      <fcntl.h>
#include      <unistd.h>
#include      <stdlib.h>
#define SHOWHOST      /* include remote machine on output */
void show_info( struct utmp* );
int main()
{
    struct utmp      current_record; /* read info into here      */
    int              utmpfd;          /* read from this descriptor */
    int              reclen = sizeof(current_record);

    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror( UTMP_FILE );      /* UTMP_FILE is in utmp.h  */
        exit(1);
    }
    while ( read(utmpfd, &current_record, reclen) == reclen )
        show_info(&current_record);
    close(utmpfd);
    return 0;                      /* went ok */
}
```


Displaying Log-In Records

```
/*
 * show_info()
 *      displays contents of the utmp struct in human readable form
 *      *note* these sizes should not be hardwired
 */
void show_info( struct utmp* utbufp);
{
    printf("%-8.8s", utbufp->ut_name);      /* the logname */
    printf(" ");                          /* a space */
    printf("%-8.8s", utbufp->ut_line);      /* the tty */
    printf(" ");                          /* a space */
    printf("%10ld", utbufp->ut_time);      /* login time */
    printf(" ");                          /* a space */
#ifdef SHOWHOST
    printf("(%s)", utbufp->ut_host);      /* the host */
#endif
    printf("\n");                          /* newline */
}
```

Compile and run it:

```
$ cc who1.c -o who1
```

```
$ ./who1
```

```
      system b  952601411 ()
      run-leve  952601411 ()
                    952601416 ()
                    952601416 ()
                    952601417 ()
                    952601417 ()
                    952601419 ()
                    952601419 ()
                    952601423 ()
                    952601566 ()
LOGIN      console  952601566 ()
            ttyt1    958240622 ()
shpyrko    ttyt2    964318862 (nas1-093.gas.swamp.org)
acotton    ttyt3    964319088 (math-guest04.williams.edu)
            ttyt4    964320298 ()
spradlin   ttyt5    963881486 (h002078c6adfb.ne.rusty.net)
dkoh       ttyt6    964314388 (128.103.223.110)
spradlin   ttyt7    964058662 (h002078c6adfb.ne.rusty.net)
king       ttyt8    964279969 (blade-runner.mit.edu)
berschba   ttyt9    964188340 (dudley.learned.edu)
rserved    ttyta    963538145 (gigue.eas.ivy.edu)
dabel      ttytb    964319455 (roam193-27.student.state.edu)
            ttytc    964319645 ()
```

Let's compare our program with the system version

```
$ who
shpyrko      ttyp2      Jul 22 22:21      (nas1-093.gas.swamp.edu)
acotton      ttyp3      Jul 22 22:24      (math-guest04.williams.edu)
spradlin     ttyp5      Jul 17 20:51      (h002078c6adfb.ne.rusty.net)
dkoh         ttyp6      Jul 22 21:06      (128.103.223.110)
spradlin     ttyp7      Jul 19 22:04      (h002078c6adfb.ne.rusty.net)
king         ttyp8      Jul 22 11:32      (blade-runner.mit.edu)
berschba     ttyp9      Jul 21 10:05      (dudley.learned.edu)
rserved      ttya       Jul 13 21:29      (gigue.eas.ivy.edu)
dabel        ttypb      Jul 22 22:30      (roam193-27.student.state.edu)
rserved      ttyd       Jul 13 21:31      (gigue.eas.harvard.edu)
dkoh         ttye       Jul 22 16:46      (128.103.223.110)
molay        ttyq0      Jul 22 20:03      (xyz73-200.harvard.edu)
cweiner      ttyq8      Jul 21 16:40      (roam175-157.student.stats.edu)
$
```

■ What We Need to Do

- Suppress blank records
- Get the log-in times correct

Writing `who2.c`

■ Suppressing blank records

○ `/usr/include/utmp.h`

```
/*      Definitions for ut_type      */
#define EMPTY      0
#define RUN_LVL    1
#define BOOT_TIME  2
#define OLD_TIME   3
#define NEW_TIME   4
#define INIT_PROCESS 5    /* Process spawned by "init" */
#define LOGIN_PROCESS 6   /* A "getty" process waiting for login */
#define USER_PROCESS 7   /* A user process */
#define DEAD_PROCESS 8    /* represents the user logged into the system.
```

○ modification

```
show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )    /* users only ! */
        return;
    printf("%-8.8s", utbufp->ut_name);          /* the username */
}
```

Writing `who2.c` (cont.)

- Displaying Log-in Time in Human-Readable Form

```
$ man -k time
```

```
$ man -k time | grep transform
```

```
$ man -k time | grep -i convert
```

- How unix stores times: `time_t`

- Unix stores times as the number of seconds since midnight, Jan 1, 1970, G.M.T. The `time_t` data type is an integer that stores a number of seconds.

- Making a `time_t` readable: `ctime`

Writing who2.c (cont.)

■ Displaying Log-in Time in Human-Readable Form

```
$ man 3 ctime
```

```
CTIME(3)
```

```
Linux Programmer's Manual
```

```
CTIME(3)
```

```
NAME
```

```
asctime, ctime, gmtime, localtime, mktime - transform  
binary date and time to ASCII
```

```
SYNOPSIS
```

```
#include <time.h>
```

```
char *asctime(const struct tm *timeptr);
```

```
char *ctime(const time_t *timep);
```

```
...
```

```
The ctime() function converts the calendar time timep into  
a string of the form
```

```
"Wed Jun 30 21:49:08 1993\n"
```

Writing who2.c (cont.)

■ Putting it All together

```
/* who2.c - read /var/adm/utmp and list info therein
 *          - suppresses empty records
 *          - formats time nicely
 */
#include      <stdio.h>
#include      <unistd.h>
#include      <utmp.h>
#include      <fcntl.h>
#include      <time.h>
#include      <stdlib.h>
/* #define SHOWHOST */

void showtime(long);
void show_info(struct utmp *);

int main()
{
    struct utmp    utbuf;          /* read info into here */
    int            utmpfd;         /* read from this descriptor */
```

✖ \$ ls -l /usr/include | more

Can I Write [who](#)?

```
    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }

    while( read(utmpfd, &utbuf, sizeof(utbuf)) == sizeof(utbuf) )
        show_info( &utbuf );
    close(utmpfd);
    return 0;
}
/*
 *      show info()
 *
 *          displays the contents of the utmp struct
 *          in human readable form
 *          * displays nothing if record has no user name
 */
void show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )
        return;

    printf("%-8.8s", utbufp->ut_name);      /* the logname */
    printf(" ");                          /* a space */
    printf("%-8.8s", utbufp->ut_line);      /* the tty */
    printf(" ");                          /* a space */
    showtime( utbufp->ut_time );           /* display time */
#ifdef SHOWHOST
    if ( utbufp->ut_host[0] != '\0' )
        printf(" (%s)", utbufp->ut_host); /* the host */
#endif
    printf("\n");                          /* newline */
}
```


Can I Write `who`?

```
void showtime( long timeval )
/*
 *   displays time in a format fit for human consumption
 *   uses ctime to build a string then picks parts out of it
 *   Note: %12.12s prints a string 12 chars wide and LIMITS
 *   it to 12chars.
 */
{
    char    *cp;                                /* to hold address of time */
    cp = ctime(&timeval);                       /* convert time to string */
                                                /* string looks like */
                                                /* Mon Feb  4 00:46:40 EST 1991 */
                                                /* 0123456789012345. */
    printf("%12.12s", cp+4 );                   /* pick 12 chars from pos 4 */
}
```

Wed Jun 30 21:49:08 1993\n

partial string!

Testing who2.c

```
$ cc who2.c -o who2
```

```
$ ./who2
```

```
rlscott  ttyp2      Jul 23 01:07
acotton  ttyp3      Jul 22 22:24
spradlin ttyp5      Jul 17 20:51
spradlin ttyp7      Jul 19 22:04
king     ttyp8      Jul 22 11:32
berschba ttyp9      Jul 21 10:05
rserverd ttypa      Jul 13 21:29
rserverd ttypd      Jul 13 21:31
molay    ttyq0      Jul 22 20:03
cweiner  ttyq8      Jul 21 16:40
mnabavi  ttyx2      Apr 10 23:11
```

```
$ who
```

```
rlscott      ttyp2      Jul 23 01:07
acotton      ttyp3      Jul 22 22:24
spradlin     ttyp5      Jul 17 20:51
spradlin     ttyp7      Jul 19 22:04
king         ttyp8      Jul 22 11:32
berschba     ttyp9      Jul 21 10:05
rserverd     ttypa      Jul 13 21:29
rserverd     ttypd      Jul 13 21:31
molay        ttyq0      Jul 22 20:03
cweiner      ttyq8      Jul 21 16:40
mnabavi      ttyx2      Apr 10 23:11
```

```
$
```

Agenda

5 min break!

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Q1: What does **cp** do?

- **cp** makes a copy of a file

\$ **cp** source-file target-file

- If there is no target file, **cp** creates it.
- If there is a target file, **cp** replaces the contents of that file with the contents of the source file.

Q2: How Does cp Create and Write?

- Creating/Truncating a File:

```
fd = create(name, 644);
```

creat		
PURPOSE	Create or zero a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = creat(char *filename, mode_t mode)	
ARGS	filename:	the name of the file
	mode:	access permission
RETURNS	-1	on error
	fd	on success

Q2: How Does **cp** Create and Write?

■ Writing to a File

```
n = write(fd, buffer, num);
```

write		
PURPOSE	Send data from memory to a file	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t result = write(int fd, void *buf, size_t amt)	
ARGS	fd	a file descriptor
	buf	an array
	amt	how many bytes to write
RETURNS	-1	on error
	num written	on success

Q3: Can I Write cp?

■ Program outline

```
open sourcefile for reading
open copyfile   for writing
+--> read from source to buffer -- eof? --+
|___ write from buffer to copy             |
                                           |
close sourcefile      <-----+
close copyfile
```




FIGURE 2.4

Copying a file by reading and writing.

Can I Write `cp`

```
/** cp1.c
 *      version 1 of cp - uses read and write with tunable buffer size
 *
 *      usage: cp1 src dest
 */
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#define BUFFERSIZE 4096
#define COPYMODE 0644

void oops(char *, char *);

main(int ac, char *av[])
{
    int    in_fd, out_fd, n_chars;
    char    buf[BUFFERSIZE];

    /* check args */
    if ( ac != 3 ){
        fprintf( stderr, "usage: %s source destination\n", *av);
        exit(1);
    }

    /* open files */
    if ( (in_fd=open(av[1], O_RDONLY)) == -1 )
        oops("Cannot open ", av[1]);

    if ( (out_fd=creat( av[2], COPYMODE)) == -1 )
        oops( "Cannot creat", av[2]);

    /* copy files */
```



```
while ( (n_chars = read(in_fd , buf, BUFFERSIZE)) > 0 )
    if ( write( out_fd, buf, n_chars ) != n_chars )
        oops("Write error to ", av[2]);
if ( n_chars == -1 )
    oops("Read error from ", av[1]);

/* close files */

if ( close(in_fd) == -1 || close(out_fd) == -1 )
    oops("Error closing files", "");
}

void oops(char *s1, char *s2)
{
    fprintf(stderr, "Error: %s ", s1);
    perror(s2);
    exit(1);
}
```

Compile and run it:

```
$ cc cp1.c -o cp1
$ ./cp1 cp1 copy.of.cp1
$ ls -l cp1 copy.of.cp1
-rw-r--r--  1 bruce  bruce      37419 Jul 23 03:12 copy.of.cp1
-rwxrwxr-x  1 bruce  bruce      37419 Jul 23 03:08 cp1
$ cmp cp1 copy.of.cp1
$
```

How well does our program respond to errors?

```
$ cp1 xxx123 file1
Error: Cannot open xxx123: No such file or directory
$ cp1 cp1 /tmp
Error: Cannot creat /tmp: Is a directory
```

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- **2.7 More Efficient File I/O: Buffering**
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

Does the Size of the Buffer Matter?

■ Example:

Ex: Filesize = 2500 bytes

If buffer = 100 bytes then
copy requires 25 read() and 25 write() calls

If buffer = 1000 bytes then
copy requires 3 read() and 3 write() calls

■ System calls consume time:

buffer size	execution time in seconds
1	50.29
4	12.81
16	3.28
64	0.96
128	0.56
256	0.37
512	0.27
1024	0.22
2048	0.19
4096	0.18
8192	0.18
16384	0.18

cp1 copying a 5MB file

Why System Calls are Time Consuming?

- It runs various kernel functions, and it also requires a shift from USER MODE to KERNEL MODE and back;
- Not only does transferring data take time but mode change takes time
- Thus, try to minimize system calls




FIGURE 2.5
Control flow during system calls.

Does This Mean That who2.c is Inefficient?

■ Yes!

- Making one system call for each line of output makes as much sense as buying pizza by the slice or eggs one at a time
- who2.c use one system call for each utmp record




■ A better idea:

- Read in a bunch of records at once
- Then, process the ones in your local storage one by one

Adding Buffering to who2.c

- We make who2.c much more efficient by using buffering to reduce system calls



File buffering with utmplib

main calls a function in `utmplib.c` to get the next struct `utmp`.

Functions in `utmplib.c` read structs 16 at a time from the disk into an array.

The kernel is called only when all 16 are used up.

FIGURE 2.6

Buffering disk data in user space.

Revised Version: who3.c

```
/* who3.c - who with buffered reads
 *          - surpresses empty records
 *          - formats time nicely
 *          - buffers input (using utmplib)
 */
```

```
#include <stdio.h>
#include <sys/types.h>
#include <utmp.h>
#include <fcntl.h>
#include <time.h>
#include <stdlib.h>
#define SHOWHOST
```

```
void show_info(struct utmp *);
void showtime(time_t);
```

```
int main()
{
```

```
    struct utmp    *utbufp,          /* holds pointer to next rec    */
                   *utmp_next();      /* returns pointer to next    */
```

```
    if ( utmp_open( UTMP_FILE ) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }
```

```
    while ( ( utbufp = utmp_next() ) != ((struct utmp *) NULL) )
        show_info( utbufp );
    utmp_close();
    return 0;
```

```
}
```

```
/*
 * show info()
```

```
...
```

```
struct utmp utbuf;
int utmpfd;

if ( ( utmpfd = open(UTMP_FILE, O_RDONLY) ) == -1 ) {
    perror( UTMP_FILE );
    exit(1);
}

while( read(utmpfd, &utbuf, sizeof(utbuf) ) == sizeof(utbuf) )
    show_info(&utbuf);
close(utmpfd);
```


■ The code for `utmplib.c`

```
/* utmplib.c - functions to buffer reads from utmp file
 *
 *      functions are
 *          utmp_open( filename )    - open file
 *          returns -1 on error
 *          utmp_next( )             - return pointer to next struct
 *          returns NULL on eof
 *          utmp_close()             - close file
 *
 *      reads NRECS per read and then doles them out from the buffer
 */
```

■ The code for utmplib.c (cont.)

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/types.h>
#include <utmp.h>

#define NRECS 16
#define NULLUT ((struct utmp *)NULL)
#define UTSIZE (sizeof(struct utmp))

static char utmpbuf[NRECS * UTSIZE];          /* storage */
static int num_recs;                          /* num stored */
static int cur_rec;                          /* next to go */
static int fd_utm = -1;                      /* read from */

utmp_open( char *filename )
{
    fd_utm = open( filename, O_RDONLY );      /* open it */
    cur_rec = num_recs = 0;                  /* no recs yet */
    return fd_utm;                          /* report */
}

struct utmp *utmp_next()
{
    struct utmp *rec;
    if ( fd_utm == -1 )                      /* error ? */
        return NULLUT;
    if ( cur_rec==num_recs && utmp_reload()==0 ) /* any more ? */
        return NULLUT;
    /* get address of next record */
    rec = ( struct utmp *) &utmpbuf[cur_rec * UTSIZE];
    cur_rec++;
    return rec;
}
```

■ The code for `utmplib.c` (cont.)

```
int utmp_reload()
/*
 *   read next bunch of records into buffer
 */
{
    int      amt_read;

                                /* read them in          */
    amt_read = read( fd_utmp , utmpbuf, NRECS * UTSIZE );

                                /* how many did we get? */
    num_recs = amt_read/UTSIZE;

                                /* reset pointer          */
    cur_rec  = 0;
    return num_recs;
}

utmp_close()
{
    if ( fd_utmp != -1 )          /* don't close if not */
        close( fd_utmp );        /* open                */
}
```

- Compile and Run

```
$ cc who3.c umtplib.c -o who3
```

```
$ ./who3
```


Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- **2.8 Buffering and the Kernel**
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

If Buffering Is So Smart, Why Doesn't the Kernel Do It?

■ It does!

- To save time, the kernel keeps copies of disk blocks in memory
- The `read()` call copies data into a process from a kernel buffer, and `write()` copies data from the process to a kernel buffer



Consequences of Kernel Buffering

- Faster “disk” I/O
- Optimized disk writes
- Need to write buffers to disk before shutdown

FIGURE 2.7

Buffering disk data in the kernel.

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- **2.9 Reading and Writing a File**
- 2.10 What to Do with System-Call Errors

Logging Out: What It Does

- The system changes a record in the utmp file.
 - Specifically what changes?

- Experiment to see how it works:
 - 1. Log in to one machine.
 - 2. Use the who1 program we wrote to see the contents of utmp.
 - 3. Log out of one of your sessions.
 - 4. Repeat 1-3 to see what happened to the utmp record.

Logging Out: How It Works

- The program that removes your name from the log has to do the following:

- **1. Open the utmp file.**

```
fd = open(UTMP_FILE, O_RDWR);
```

- **2. Read the utmp file until it finds the record for your terminal.**

```
While( read(fd, rec, utmplen) == utmplen ) /* get next record */  
    If( strcmp(rec.ut_line, myline) == 0) /* what, my line? */  
        revise_entiry(); /* remove my name */
```

Logging Out: How It Works (cont.)

- The program that removes your name from the log has to do the following:

- **3. Write a revised utmp record in its place.**

- How do we write the revised record back to the file?
 - › If we just call `write`, our code updates the next record

Q: How can a program change the current read-write position in a file?
A: The `lseek` system call.

- **4. Close the utmp file.**

```
close(fd);
```

Moving the Current Position: `lseek()`

- Every open file has a current position
 - The current position belongs to the connection to the file, not to the file

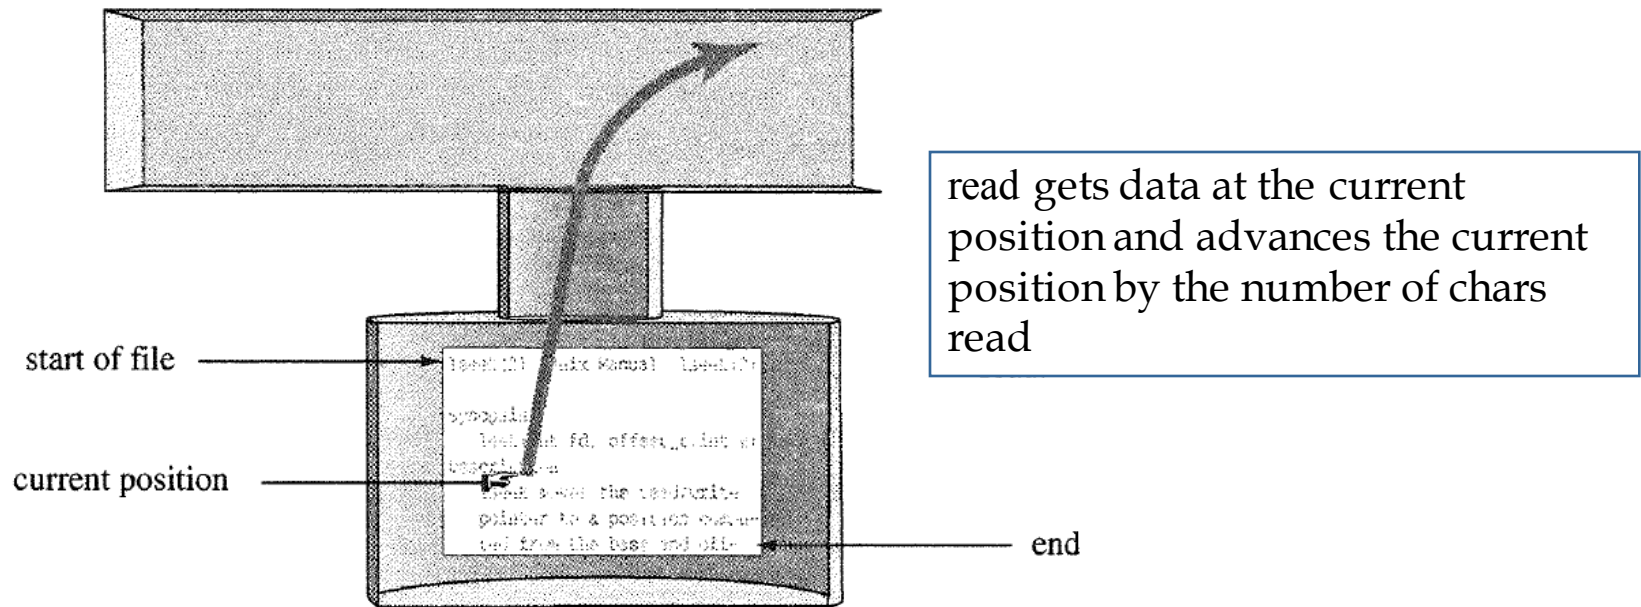


FIGURE 2.8

Every open file has a current position.

Moving the Current Position: `lseek()`

- The `lseek()` system call lets you change the current position of an open file:

```
lseek(fd, -(sizeof(struct utmp)), SEEK_CUR);  
lseek(fd, 10 * sizeof(struct utmp), SEEK_SET);
```

```
lseek(fd, 0, SEEK_END);  
write(fd, "hello", strlen("hello"));
```

```
lseek(fd, 0, SEEK_CUR)
```

Moving the Current Position: **lseek()**

lseek		
PURPOSE	Set file pointer to specified offset in file	
INCLUDE	#include <sys/types.h> #include <unistd.h>	
USAGE	off_t oldpos = lseek(int fd, off_t dist, int base)	
ARGS	fd:	file descriptor
	dist:	a distance in bytes
	base:	SEEK_SET => start of file SEEK_CUR => current position SEEK_END => end of file
RETURNS	-1	on error
	or	the previous position in the file

Code to Log Out from a Terminal

```
/*
 * logout_tty(char *line)
 *   marks a utmp record as logged out
 *   does not blank username or remote host
 *   returns -1 on error, 0 on success
 */
int logout_tty(char *line)
{
    int      fd;
    struct utmp rec;
    int      len = sizeof(struct utmp);
    int      retval = -1 ;                /* pessimism */

    if ( (fd = open(UTMP_FILE,O_RDWR)) == -1 ) /* open file */
        return -1;

    /* search and replace */
    while ( read(fd, &rec, len) == len)
        if ( strcmp(rec.ut_line, line, sizeof(rec.ut_line)) == 0)
        {
            rec.ut_type = DEAD_PROCESS;    /* set type */
            if ( time( &rec.ut_time ) != -1 ) /* and time */
                if ( lseek(fd, -len, SEEK_CUR) != -1 ) /* back up */
                    if ( write(fd, &rec, len) == len ) /* update */
                        retval = 0;                /* success! */
            break;
        }

    /* close the file */
    if ( close(fd) == -1 )
        retval = -1;
    return retval;
}
```

Agenda

- 2.2 Asking about who
- 2.3 What Does who Do?
- 2.4 How Does who Do It?
- 2.5 Can I Write who?
- 2.6 Writing cp (read and write)
- 2.7 More Efficient File I/O: Buffering
- 2.8 Buffering and the Kernel
- 2.9 Reading and Writing a File
- 2.10 What to Do with System-Call Errors

What to Do with System-Call Errors

- When open cannot open a file, it returns -1
- When read cannot read data, it returns -1
- When lseek cannot seek, it returns -1
- System calls return -1 when something goes wrong
- Your programs should test the return value of every system call they make and take intelligent action when errors occur

How to Identify What Went Wrong: `errno`

- The kernel tells your program the cause of the error by storing an error code in a global variable called `errno`
- The manpage for `errno(3)` and the file `<errno.h>` include the error-code symbols and numeric codes

```
$ man 3 errno
```

```
#define EPERM          1      /* Operation not permitted */
#define ENOENT         2      /* No such file or directory */
#define ESRCH          3      /* No such process */
#define EINTR          4      /* Interrupted system call */
#define EIO            5      /* I/O error */
```

Different Responses to Different Errors

```
#include <errno.h>

extern int errno;

int sample()
{
    int fd;
    fd = open("file", O_RDONLY);
    if ( fd == -1 )
    {
        printf("Cannot open file: ");
        if ( errno == ENOENT )
            printf("There is no such file.");
        else if ( errno == EINTR )
            printf("Interrupted while opening file.");

        else if ( errno == EACCESS )
            printf("You do not have permission to open file.");
        ...
    }
}
```

Reporting Errors: perror(3)

- If you want to print a message describing the error, you could test the value of `errno` and print different messages for different values.
- Print a system error message

```
int sample()
{
    int fd;
    fd = open("file", O_RDONLY);
    if ( fd == -1 )
    {
        perror("Cannot open file");
        return;
    }
    ...
}
```

Objectives

■ Ideas and Skills

- The role and use of on-line documentation ([man](#))
- The Unix file interface: [open](#), [read](#), [write](#), [lseek](#), [close](#)
- [Reading, creating, and writing files](#)
- [File descriptors](#)
- [Buffering](#): user level and kernel level
- Kernel mode, user mode, and the [cost of system calls](#)
- How Unix represents time, [how to format Unix time](#)
- Using the [utmp](#) file to find list of current users
- [Detecting and reporting errors](#) in system calls

■ System Calls and Functions

- [open](#), [read](#), [write](#), [create](#), [lseek](#), [close](#)
- [perror](#)

■ UNIX/LINUX Commands

- [man](#), [who](#), [cp](#), [login](#)

Next Time

- Chapter 3 : Directories and File Properties