

Directories and File Properties

Prof. Seokin Hong

Kyungpook National University

Fall 2018

Objectives

■ Ideas and Skills

- A directory is a list of files
- How to read a directory
- Types of files and how to determine the type of a file
- Properties of files and how to determine properties of a file
- Bit sets and bit masks
- User and group ID numbers

■ System Calls and Functions

- opendir, readdir, closedir, seekdir
- stat
- chmod, chown, utime
- rename

■ Commands

- ls

Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How Dos Is Work?
- Can I Write Is?
- Writing Is -l
- Three Special Bits
- Setting and Modifying the Properties of a File

ls

- lists names of files and reports file attributes
- Example:

```
$ ls
Makefile      docs          ls2.c         s.tar         statdemo.c    tail1.c
chap03        ls1.c         old_src       stat1.c       tail1
$
```

```
$ ls -l
total 108
-rw-rw-r--  2 bruce  users      345 Jul 29 11:05 Makefile
-rw-rw-r--  1 bruce  users    27521 Aug  1 12:14 chap03
drwxrwxr-x  2 bruce  users     1024 Aug  1 12:15 docs
Type&permission  links  owner    group      size    modified-date/time  name
```

ls

- Listing other directories and reporting on other files

| Asking ls about Other Directories and Their Files | |
|---|---|
| Example | Action |
| <code>ls /tmp</code> | list names of files in <code>/tmp</code> directory |
| <code>ls -l docs</code> | show attributes of files in <code>docs</code> directory |
| <code>ls -l ../Makefile</code> | show attributes of <code>../Makefile</code> |
| <code>ls *.c</code> | list names of files matching pattern <code>*.c</code> |

Popular Command-Line Options

■ Options:

| Command | Action |
|---------------------|------------------------|
| <code>ls -a</code> | shows <u>“.”-files</u> |
| <code>ls -lu</code> | shows last-read time |
| <code>ls -s</code> | shows size in blocks |
| <code>ls -t</code> | sorts in time order |
| <code>ls -F</code> | shows file types |

`ls -al`

■ A remark on Dot-Files (hidden file)

- **ls** does not list the name of a file if the first character of the filename is a dot.
- Some programs use dot filenames in a user's home directory to store user preferences.

```
root@DESKTOP-K4MA2V5:~# ls -l
```

합계 24

```
drwxrwxrwx 0 root root 512 2월 14 10:24 adir
-rw-rw-rw- 1 root root 27 2월 28 14:39 cat.test
-rw-rw-rw- 1 root root 9525 2월 27 16:46 etc.listing
-rwxrwxrwx 1 root root 8600 3월 12 19:06 hello
-rw-rw-rw- 1 root root 61 3월 12 19:06 hello.c
-rw-rw-rw- 1 root root 33 2월 21 15:05 test.c
-rw-rw-rw- 1 root root 0 2월 12 15:02 userlist
-rw-rw-rw- 1 root root 13 3월 8 10:09 vitest.txt
```

```
root@DESKTOP-K4MA2V5:~# ls -al
```

합계 32

```
drwx----- 0 root root 512 3월 12 19:06 .
drwxr-xr-x 0 root root 512 8월 25 2017 ..
-rw----- 1 root root 96 2월 28 10:45 .bash_history
-rw-r--r-- 1 root root 3106 10월 23 2015 .bashrc
drwxrwxrwx 0 root root 512 3월 12 18:59 .nano
-rw-r--r-- 1 root root 148 8월 18 2015 .profile
drwxr-xr-x 0 root root 512 3월 12 18:58 .vim
-rw----- 1 root root 3370 3월 12 18:59 .viminfo
drwxrwxrwx 0 root root 512 2월 14 10:24 adir
-rw-rw-rw- 1 root root 27 2월 28 14:39 cat.test
-rw-rw-rw- 1 root root 9525 2월 27 16:46 etc.listing
-rwxrwxrwx 1 root root 8600 3월 12 19:06 hello
-rw-rw-rw- 1 root root 61 3월 12 19:06 hello.c
-rw-rw-rw- 1 root root 33 2월 21 15:05 test.c
-rw-rw-rw- 1 root root 0 2월 12 15:02 userlist
-rw-rw-rw- 1 root root 13 3월 8 10:09 vitest.txt
```

```
root@DESKTOP-K4MA2V5:~#
```

- Dot-file (hidden file)
- Dot-files at home directory are typically used for user preferences of programs

So, what does **ls** do?

- **ls** does two things

1. Lists the contents of directories
2. Displays information about files

- We need to learn:

1. How to list the contents of a directory
2. How to obtain and display properties of a file
3. How to determine if a name refers to a file or a directory

Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

File System Tree

- The **disk** is organized as a **tree of directories**, each of which contains files or directories.
- The commands **cd**, **pwd**, **ls** allow us to explore a file system




FIGURE 3.1

A tree of directories.

How Dos **ls** Work?

■ Outline :

```
        open directory
+--> read entry          -end of dir?--+
|__ display file info    |
    close directory      <-----+

```

■ It looks just like the logic for who!

■ Difference?

- The main difference is that the `who` opens and reads *from a file* (*utmp*, *wtmp*)
- `ls` opens and reads its data *from a directory*.

What is a Directory?

- A directory is a kind of file that contains a list of names of files and directories.
- Unlike a regular file, a directory never empty
 - Every directory contains two specific items: **.** **and** **..**
 - **dot(.)** is the name of the current directory,
 - **dotdot(..)** is the name of the directory one level up.

Do open, read, and close work for directories?

- Answer 1: on old versions of Unix, that was the only way
 - On some versions of Unix, you still can, but not for all directories

- Answer 2: It is a bad idea to use open, read and close to list contents of directory. Why?
 - Unix allows various disk formats to appear as part of a single tree.
 - It supports Mac HFS, FAT, FAT32, lots of Unix flavors;
 - Thus, using **read** to process each type would require knowing the format of the records for each type of directory

How do I read a Directory?

```
$ man -k direct
```

```
$ man -k direct | grep read
```

```
DXmHelpSystemDisplay (3X) - Displays a topic or directory of the  
help file in Bookreader.
```

```
opendir, readdir, readdir_r, telldir, seekdir, rewinddir, closedir (3) -  
Performs operations on directories
```

```
$ man 3 readdir
```

```
opendir(3) opendir(3)
```

```
NAME
```

```
opendir, readdir, readdir_r, telldir, seekdir, rewinddir, closedir -  
Performs operations on directories
```

LIBRARY

Standard C Library (libc.a)

SYNOPSIS

```
#include <sys/types.h>
#include <dirent.h>
```

```
DIR *opendir (
    const char *dir_name );
```

```
struct dirent *readdir (
    DIR *dir_pointer );
```

```
int readdir_r (
    DIR *dir_pointer,
    struct dirent *entry,
    struct dirent **result);
```

```
long telldir (
    DIR *dir_pointer );
```

```
void seekdir (
    DIR *dir_pointer,
    long location );
```

```
void rewinddir (
    DIR *dir_pointer );
```

```
int closedir (
    DIR *dir_pointer );
```

[more] (11%)

```
#include <dirent.h>
```

```
opendir(char *)  
  creates a connection,  
  returns a DIR *
```

```
readdir(DIR *)  
  reads next record,  
  returns a pointer  
  to a struct dirent
```

```
closedir(DIR *)  
  closes a connection
```




FIGURE 3.2

Reading entries from a directory.

Reading the contents of a directory

- We read the entire directory by calling `readdir()`
- Each `readdir()` call returns a pointer to the next record, a variable of type `struct dirent`

```
struct dirent *readdir (  
    DIR *dir_pointer );
```

NAME

dirent - file system independent directory entry

SYNOPSIS

```
#include <dirent.h>
```

DESCRIPTION

Different file system types may have different directory entries. The `dirent` structure defines a file system independent directory entry, which contains information common to directory entries in different file system types. A set of these structures is returned by the `getdents(2)` system call.

The `dirent` structure is defined:

```
struct dirent {
    ino_t      d_ino;
    off_t      d_off;
    unsigned short d_reclen;
    char        d_name[1];
};
```

Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

Writing ls1.c

- Logic for listing a directory:

```
main()
    opendir
    while ( readdir )
        print d_name
    closedir
```

```

/** ls1.c
**  purpose - list contents of directory or directories
**  action - if no args, use .  else list files in args
**/
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>

void do_ls(char []);

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}

void do_ls( char dirname[] )
/*
 *   list files in directory called dirname
 */
{
    DIR          *dir_ptr;           /* the directory */
    struct dirent *direntp;          /* each entry   */

    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            printf("%s\n", direntp->d_name );
        closedir(dir_ptr);
    }
}

```

- Compile and run it:

```
$ cc -o ls1 ls1.c
```

```
$ ls1
```

```
.  
..  
s.tar  
tail1  
Makefile  
ls1.c  
ls2.c  
chap03  
old_src  
docs  
ls1  
stat1.c  
statdemo.c  
tail1.c
```

```
$ ls
```

```
Makefile      docs          ls1.c         old_src       stat1.c       tail1  
chap03        ls1           ls2.c         s.tar         statdemo.c   tail1.c  
$
```

✖

```
$ ./ls1
```

```
$ ./ls1 . /tmp /usr
```

Agenda

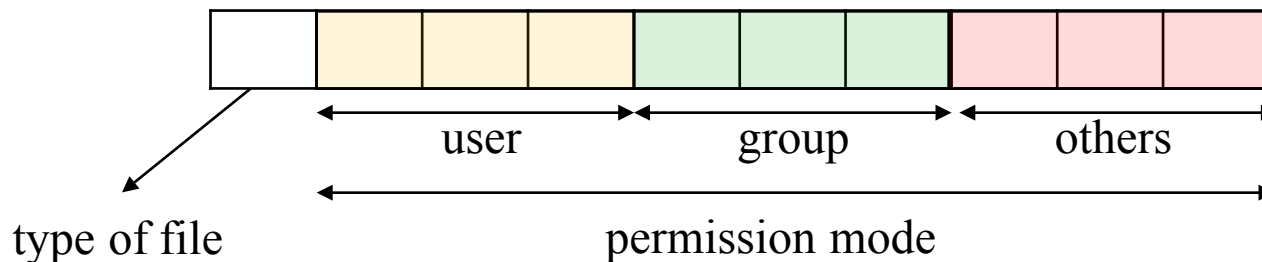
- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

What Does `ls -l` Do?

- **ls** does two different types of things
 - lists directories and files
 - display information about directories and files

```
$ ls -l
total 108
-rw-rw-r-- 2 bruce  users      345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce  users    27521 Aug  1 12:14 chap03
drwxrwxr-x 2 bruce  users     1024 Aug  1 12:15 docs
```

type and permission links owner group size last-modified time name



- : regular file
d : directory

How Does ls -l work?

- How can we get information (status/properties) about a file?
 - **stat** system call is used to retrieve file status
 - Search the manual:
\$ man -k file | grep -i status

```
root@DESKTOP-K4MA2V5:~# man -k file | grep status
fileno (3)          - check and reset stream status
fstat (2)           - get file status
fstat64 (2)         - get file status
fstatat (2)         - get file status
fstatat64 (2)       - get file status
lstat (2)           - get file status
lstat64 (2)         - get file status
newfstatat (2)      - get file status
oldfstat (2)        - get file status
oldlstat (2)        - get file status
oldstat (2)         - get file status
stat (1)            - display file or file system status
stat (2)            - get file status
stat64 (2)          - get file status
root@DESKTOP-K4MA2V5:~#
```

- \$ man 2 stat

```
root@DESKTOP-K4MA2V5: ~  
STAT(2)  
Programmer's Manual  
STAT(2)  
  
NAME  
    stat, fstat, lstat, fstatat - get file status  
  
SYNOPSIS  
    #include <sys/types.h>  
    #include <sys/stat.h>  
    #include <unistd.h>  
  
    int stat(const char *pathname, struct stat *buf);  
    int fstat(int fd, struct stat *buf);  
    int lstat(const char *pathname, struct stat *buf);  
  
    #include <fcntl.h>          /* Definition of AT_* constants */  
    #include <sys/stat.h>  
  
    int fstatat(int dirfd, const char *pathname, struct stat *buf,  
                int flags);  
  
Feature Test Macro Requirements for glibc (see feature_test_macros(7)):  
Manual page stat(2) line 1 (press h for help or q to quit)
```

All of these system calls return a stat structure, which contains the following fields:

```
struct stat {  
    dev_t    st_dev;        /* ID of device containing file */  
    ino_t    st_ino;        /* inode number */  
    mode_t    st_mode;      /* protection */  
    nlink_t   st_nlink;     /* number of hard links */  
    uid_t    st_uid;        /* user ID of owner */  
    gid_t    st_gid;        /* group ID of owner */  
    dev_t    st_rdev;       /* device ID (if special file) */  
    off_t     st_size;      /* total size, in bytes */  
    blksize_t st_blksize;   /* blocksize for filesystem I/O */  
    blkcnt_t  st_blocks;    /* number of 512B blocks allocated */  
  
    /* Since Linux 2.6, the kernel supports nanosecond  
       precision for the following timestamp fields.  
       For the details before Linux 2.6, see NOTES. */  
  
    struct timespec st_atim; /* time of last access */  
    struct timespec st_mtim; /* time of last modification */  
    struct timespec st_ctim; /* time of last status change */  
};
```

Ans: The **stat** system call gets file information

■ How does **stat** work:

- A file is stored on the disk, and it has contents and a set of attributes (e.g., size, owner ID, etc)
- The process defines a buffer of type **struct state**
- And then asks the kernel to copy file information from the disk to the buffer




FIGURE 3.3

Reading file properties using `stat`.

stat

PUPPOSE Obtain information about a file

INCLUDE #include <sys/stat.h>

USAGE int result = stat(char *fname, struct stat *bufp)

AGRS fname name of file
 bufp pointer to buffer

RETURNS -1 if error
 0 if success

```
/* filesize.c - prints size of passwd file */
#include <stdio.h>
#include <sys/stat.h>

int main()
{
    struct stat infobuf;                /* place to store info */
    if ( stat( "/etc/passwd", &infobuf) == -1 ) /* get info */
        perror("/etc/passwd");
    else
        printf(" The size of /etc/passwd is %d\n", infobuf.st_size );
}
```

What other information does **stat** provide?

- Members of struct stat

- Described in /usr/include/sys/stat.h

| | |
|----------|------------------------------|
| st_mode | type and permission |
| st_uid | ID of owner |
| st_gid | ID of group |
| st_size | number of bytes in file |
| st_nlink | number of links to file |
| st_mtime | last content-modified time |
| st_atime | last-accessed time |
| st_ctime | last properties-changed time |

-
- Use stat to get file info for that name
 - Display the items in the struct

```
/* fileinfo.c - use stat() to obtain and print file properties
 *             - some members are just numbers...
 */
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>

void show_stat_info(char *, struct stat *);

int main(int ac, char *av[])
{
    struct stat info;          /* buffer for file info */

    if (ac>1)
        if( stat(av[1], &info) != -1 ){
            show_stat_info( av[1], &info );
            return 0;
        }
        else
            perror(av[1]); /* report stat() errors */
    return 1;
}
```

```
void show_stat_info(char *fname, struct stat *buf)
/*
 * displays some info from stat in a name=value format
 */
{
    printf("    mode: %o\n", buf->st_mode);          /* type + mode */
    printf("    links: %d\n", buf->st_nlink);        /* # links      */
    printf("    user: %d\n", buf->st_uid);           /* user id      */
    printf("    group: %d\n", buf->st_gid);          /* group id     */
    printf("    size: %d\n", buf->st_size);          /* file size    */
    printf("modtime: %d\n", buf->st_mtime);         /* modified     */
    printf("    name: %s\n", fname );              /* filename     */
}
```

-
- Compile and run it :

```
$ cc -o fileinfo fileinfo.c
```

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664
```

```
links: 1
```

```
user: 500
```

```
group: 120
```

```
size: 1106
```

```
modtime: 965158604
```


```
name: fileinfo.c
```

```
$ ls -l fileinfo.c
```

```
-rw-rw-r--  1 bruce    users          1106 Aug  1 15:36 fileinfo.c
```

Converting file mode to a string

- File type and permission bits are stored in the **st_mode** member



- **Type:** file types

- 4 bits means 16 possible patterns.
- Each pattern can correspond to a file type.

- **Permission bits :**

- Access permission of user, group, others for the file
- 1 indicates the permission is granted
- 0 indicates the permission is denied

```
$ ./fileinfo fileinfo.c
mode: 100664
```

How to read subfields: Masking

- How do we examine a bit or sub-field?

- ex) 100664 (base 8) → -rw-rw-r--

- Use “bitwise AND (&)” to MASK

| | | | | | | | | | | | | | | | | | |
|-------|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|--------|
| ○ Ex) | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | value |
| & | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | mask |
| = | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | result |

FIGURE 3.6

Applying a bitmask.

Using Masking to decode permission bits

■ 100664 (base 8) → -rw-rw-r--

```
/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and
 * nine letters that correspond to the bits
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" ); /* default=no perms */
    if ( S_ISDIR(mode) ) str[0] = 'd'; /* directory? */
    if ( S_ISCHR(mode) ) str[0] = 'c'; /* char devices */
    if ( S_ISBLK(mode) ) str[0] = 'b'; /* block device */

    if ( mode & S_IRUSR ) str[1] = 'r'; /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r'; /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

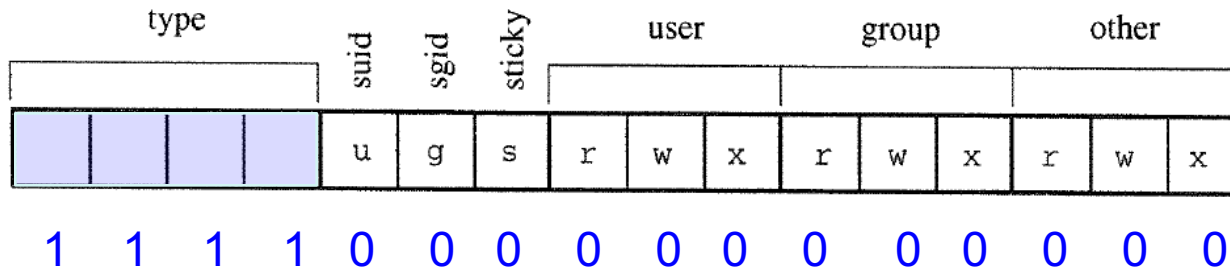
    if ( mode & S_IROTH ) str[7] = 'r'; /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

| | | |
|---------|-------|---|
| S_IRWXU | 00700 | owner has read, write, and execute permission |
| S_IRUSR | 00400 | owner has read permission |
| S_IWUSR | 00200 | owner has write permission |
| S_IXUSR | 00100 | owner has execute permission |
| S_IRWXG | 00070 | group has read, write, and execute permission |
| S_IRGRP | 00040 | group has read permission |
| S_IWGRP | 00020 | group has write permission |
| S_IXGRP | 00010 | group has execute permission |

masks defined in <sys/stat.h>

Using Masking to decode file types

- Mask and file types defined in <sys/stat.h>



```
#define S_IFMT      0170000    /* type of file */
#define S_IFREG     0100000    /* regular */
#define S_IFDIR     0040000    /* directory */
#define S_IFBLK     0060000    /* block special */
#define S_IFCHR     0020000    /* character special */
#define S_IFIFO     0010000    /* fifo */
#define S_IFLNK     0120000    /* symbolic link */
#define S_IFSOCK    0140000    /* socket */
```

mask

File types

※ octal

```
if ( (info.st_mode & 0170000) == 0040000 )
    printf("this is a directory.");
```

File types

```
drwxr-xr-x 2 root root    0 Jan  1  1970 home
```

- **Regular** : regular file, marked with **-**
- **Directory** : directory. marked with **d**
- **Symbolic link** : a reference to another file. marked with **l**
- **Socket** : file used for inter-process communication that enable packetized-communication between two processes. communication can extend beyond localhost. marked with an **s**
- **Block special** : interface that allows an application to interact with a hardware devices. It provides buffered access to the hardware. marked with **b**
- **Character special** : interface that allows an application to interact with a hardware devices. It provides un-buffered, direct access to the hardware. marked with **c**
- **FIFO (named pipe)** : file used for inter-process communication within a host. marked with **p**

- **Macros** defined in `<sys/stat.h>`

```
/*
 *      File type macros
 */

#define S_ISFIFO(m)      (((m)&(0170000)) == (0010000))
#define S_ISDIR(m)       (((m)&(0170000)) == (0040000))
#define S_ISCHR(m)       (((m)&(0170000)) == (0020000))
#define S_ISBLK(m)       (((m)&(0170000)) == (0060000))
#define S_ISREG(m)       (((m)&(0170000)) == (0100000))

if ( S_ISDIR(info.st_mode) )
    printf("this is a directory.");
```

Converting **User ID** to Strings

```
$ ./fileinfo fileinfo.c
  mode: 100664
  links: 1
  user: 500
  group: 120
  size: 1106
modtime: 965158604
  name: fileinfo.c
$ ls -l fileinfo.c
-rw-rw-r-- 1 bruce  users      1106 Aug  1 15:36 fileinfo.c
```

Converting **User ID** to Strings

- Library function `getpwuid()` Provides Access to the Complete List of Users
 - Defined in `/usr/include/pwd.h`

```
root@DESKTOP-K4MA2V5: ~  
GETPWNAM(3) Linux Programmer's Manual  
NAME  
    getpwnam, getpwnam_r, getpwuid, getpwuid_r - get password file entry  
SYNOPSIS  
    #include <sys/types.h>  
    #include <pwd.h>  
  
    struct passwd *getpwnam(const char *name);  
    struct passwd *getpwuid(uid_t uid);
```

Converting User ID to Strings

```
struct passwd *getpwuid(uid_t uid);
```

※ \$ man getpwuid

```
/* The passwd structure. */
struct passwd
{
    char *pw_name;           /* Username. */
    char *pw_passwd;         /* Password. */
    __uid_t pw_uid;          /* User ID. */
    __gid_t pw_gid;          /* Group ID. */
    char *pw_gecos;          /* Real name. */
    char *pw_dir;            /* Home directory. */
    char *pw_shell;          /* Shell program. */
};
```

```
/*
 * returns a username associated with the specified uid
 * NOTE: does not work if there is no username
 */
char *uid_to_name( uid_t uid )
{
    return getpwuid(uid)->pw_name ;
}
```

Converting Group ID to Strings

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664
```

```
links: 1
```

```
user: 500
```

```
group: 120
```

```
size: 1106
```

```
modtime: 965158604
```

```
name: fileinfo.c
```

```
$ ls -l fileinfo.c
```

```
-rw-rw-r--  1 bruce  users      1106 Aug  1 15:36 fileinfo.c
```

Converting Group ID to Strings

- `getgrgid()` provides access to the list of groups
 - Defined in `/usr/include/grp.h`

```
struct group *getgrgid(gid_t gid);    ※ $ man getgrgid
```

```
struct group {
    char    *gr_name;           /* group name */
    char    *gr_passwd;         /* group password */
    gid_t   gr_gid;             /* group ID */
    char    **gr_mem;           /* group members */
};
```

```
/*
 * returns a groupname associated with the specified gid
 * NOTE: does not work if there is no groupname
 */
char *gid_to_name( gid_t gid )
{
    return getgrgid(gid)->gr_name ;
}
```

Putting It All Together: ls2.c

```
$ cc -o ls1 ls1.c
```

```
$ ls1
```

```
.  
..  
s.tar  
tail1  
Makefile  
ls1.c  
ls2.c  
chap03  
old_src  
docs  
ls1  
stat1.c  
statdemo.c  
tail1.c
```

```
$ cc -o fileinfo fileinfo.c
```

```
$ ./fileinfo fileinfo.c
```

```
mode: 100664  
links: 1  
user: 500  
group: 120  
size: 1106  
modtime: 965158604  
name: fileinfo.c
```

```
$ ls2
```

| | | | | | | | | |
|------------|---|-------|-------|-------|-----|----|-------|----------|
| drwxrwxr-x | 4 | bruce | bruce | 1024 | Aug | 2 | 18:18 | . |
| drwxrwxr-x | 5 | bruce | bruce | 1024 | Aug | 2 | 18:14 | .. |
| -rw-rw-r-- | 1 | bruce | users | 30720 | Aug | 1 | 12:05 | s.tar |
| -rwxrwxr-x | 1 | bruce | users | 37351 | Aug | 1 | 12:13 | tail1 |
| -rw-rw-r-- | 2 | bruce | users | 345 | Jul | 29 | 11:05 | Makefile |
| -rw-r--r-- | 1 | bruce | users | 723 | Aug | 1 | 14:26 | ls1.c |
| -rw-r--r-- | 1 | bruce | users | 3045 | Feb | 15 | 03:51 | ls2.c |
| -rw-rw-r-- | 1 | bruce | users | 27521 | Aug | 1 | 12:14 | chap03 |
| drwxrwxr-x | 2 | bruce | users | 1024 | Aug | 1 | 12:14 | old_src |
| drwxrwxr-x | 2 | bruce | users | 1024 | Aug | 1 | 12:15 | docs |
| -rwxrwxr-x | 1 | bruce | bruce | 37048 | Aug | 1 | 14:26 | ls1 |

```
/* ls2.c
 *      purpose  list contents of directory or directories
 *      action   if no args, use .  else list files in args
 *      note     uses stat and pwd.h and grp.h
 *      BUG: try ls2 /tmp
 */
#include      <stdio.h>
#include      <sys/types.h>
#include      <dirent.h>
#include      <sys/stat.h>
#include      <string.h>

void do_ls(char[]);
void dostat(char *);
void show_file_info( char *, struct stat *);
void mode_to_letters( int , char [] );
char *uid_to_name( uid_t );
char *gid_to_name( gid_t );

main(int ac, char *av[])
{
    if ( ac == 1 )
        do_ls( "." );
    else
        while ( --ac ){
            printf("%s:\n", *++av );
            do_ls( *av );
        }
}
```

```
void do_ls( char dirname[] )
/*
 *      list files in directory called dirname
 */
{
    DIR          *dir_ptr;          /* the directory */
    struct dirent *direntp;          /* each entry */
    if ( ( dir_ptr = opendir( dirname ) ) == NULL )
        fprintf(stderr, "ls1: cannot open %s\n", dirname);
    else
    {
        while ( ( direntp = readdir( dir_ptr ) ) != NULL )
            dostat( direntp->d_name );
        closedir(dir_ptr);
    }
}

void dostat( char *filename )
{
    struct stat info;
    if ( stat(filename, &info) == -1 )          /* cannot stat */
        perror( filename );                    /* say why */
    else                                          /* else show info */
        show_file_info( filename, &info );
}
```

```
void show_file_info( char *filename, struct stat *info_p )
/*
 * display the info about 'filename'. The info is stored in struct at
 *info_p
 */
{
    char    *uid_to_name(), *ctime(), *gid_to_name(), *filemode();
    void    mode_to_letters();
    char    modestr[11];

    mode_to_letters( info_p->st_mode, modestr );

    printf( "%s"      , modestr );
    printf( "%4d "    , (int) info_p->st_nlink);
    printf( "%-8s "   , uid_to_name(info_p->st_uid) );
    printf( "%-8s "   , gid_to_name(info_p->st_gid) );
    printf( "%8ld "   , (long)info_p->st_size);
    printf( "%.12s "  , 4+ctime(&info_p->st_mtime));
    printf( "%s\n"    , filename );
}
```

```
/*
 * utility functions
 */

/*
 * This function takes a mode value and a char array
 * and puts into the char array the file type and the
 * nine letters that correspond to the bits in mode.
 * NOTE: It does not code setuid, setgid, and sticky
 * codes
 */
void mode_to_letters( int mode, char str[] )
{
    strcpy( str, "-----" );          /* default=no perms */

    if ( S_ISDIR(mode) ) str[0] = 'd';  /* directory?      */
    if ( S_ISCHR(mode) ) str[0] = 'c';  /* char devices    */
    if ( S_ISBLK(mode) ) str[0] = 'b';  /* block device    */

    if ( mode & S_IRUSR ) str[1] = 'r';  /* 3 bits for user */
    if ( mode & S_IWUSR ) str[2] = 'w';
    if ( mode & S_IXUSR ) str[3] = 'x';

    if ( mode & S_IRGRP ) str[4] = 'r';  /* 3 bits for group */
    if ( mode & S_IWGRP ) str[5] = 'w';
    if ( mode & S_IXGRP ) str[6] = 'x';

    if ( mode & S_IROTH ) str[7] = 'r';  /* 3 bits for other */
    if ( mode & S_IWOTH ) str[8] = 'w';
    if ( mode & S_IXOTH ) str[9] = 'x';
}
```

```
#include <pwd.h>
```

```
char *uid_to_name( uid_t uid )
/*
 *      returns pointer to username associated with uid, uses getpw()
 */
{
    struct passwd *getpwuid(), *pw_ptr;
    static char numstr[10];
    if ( ( pw_ptr = getpwuid( uid ) ) == NULL ){
        sprintf(numstr,"%d", uid);
        return numstr;
    }
    else
        return pw_ptr->pw_name ;
}
```

```
#include <grp.h>
```

```
char *gid_to_name( gid_t gid )
/*
 *      returns pointer to group number gid. used getgrgid(3)
 */
{
    struct group *getgrgid(), *grp_ptr;
    static char numstr[10];
    if ( ( grp_ptr = getgrgid(gid) ) == NULL ){
        sprintf(numstr,"%d", gid);
        return numstr;
    }
    else
        return grp_ptr->gr_name;
}
```

Result

```
$ ./ls2
drwxrwxr-x 4 bruce bruce 1024 Aug 2 18:18 .
drwxrwxr-x 5 bruce bruce 1024 Aug 2 18:14 ..
-rw-rw-r-- 1 bruce users 30720 Aug 1 12:05 s.tar
-rwxrwxr-x 1 bruce users 37351 Aug 1 12:13 tail1
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-r--r-- 1 bruce users 723 Aug 1 14:26 ls1.c
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
```

```
$ ls -l
total 189
-rw-rw-r-- 2 bruce users 345 Jul 29 11:05 Makefile
-rw-rw-r-- 1 bruce users 27521 Aug 1 12:14 chap03
drwxrwxr-x 2 bruce users 1024 Aug 1 12:15 docs
-rwxrwxr-x 1 bruce bruce 37048 Aug 1 14:26 ls1
-rw-r--r-- 1 bruce users 723 Aug 1 14:26 ls1.c
-rwxrwxr-x 2 bruce bruce 42295 Aug 2 18:18 ls2
-rw-r--r-- 1 bruce users 3045 Feb 15 03:51 ls2.c
```

} sorting




Summary

```
readdir() struct dirent {
    ino_t      d_ino;      /* inode number */
    off_t      d_off;      /* not an offset; see NOTES */
    unsigned short d_reclen; /* length of this record */
    unsigned char d_type;   /* type of file; not supported
                             by all filesystem types */
    char        d_name[256]; /* filename */
};
```

```
stat() struct stat {
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t      st_mode;    /* protection */
    nlink_t     st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t   st_blksize; /* blocksize for filesystem I/O */
    blkcnt_t    st_blocks;  /* number of 512B blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

```
getpwuid( struct passwd {
    char *pw_name;      /* username */
    char *pw_passwd;    /* user password */
    uid_t pw_uid;       /* user ID */
    gid_t pw_gid;       /* group ID */
    char *pw_gecos;     /* user information */
    char *pw_dir;       /* home directory */
    char *pw_shell;     /* shell program */
};
```

```
getgrgid( struct group {
    char *gr_name;      /* group name */
    char *gr_passwd;    /* group password */
    gid_t gr_gid;       /* group ID */
    char **gr_mem;      /* group members */
};
```




Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- **Three Special Bits**
- Setting and Modifying the Properties of a File

The Three Special Bits

- The `st_mode` member of the `stat` structure:



- Three special bits are used to activate special properties of a file
 - `suid`(set-user-ID) bit
 - `sgid`(set-group-ID) bit
 - sticky bit

1. The Set-User-ID Bit

- How can a regular user change his or her password?
 - Use the passwd command!
 - But, how does the passwd command work?

```
$ ls -l /etc/passwd
-rw-r--r--  1 root    root          894 Jun 20 19:17 /etc/passwd
```

Problem:

Changing your password means changing your record in the file `/etc/passwd`, but you do **NOT** have **permission** to write to that file.

Only the user named **root** has write permission.

1. The Set-User-ID Bit

- Solution: Give permission to the program, not to you.

```
$ ls -l /usr/bin/passwd  
-r-sr-xr-x  1 root      bin           15725 Oct 31  1997 /usr/bin/passwd
```

- The program you use to change your password, /usr/bin/passwd or /bin/passwd, is owned by root and has the set-user-ID (SUID) bit set.
- That **SUID** bit tells the kernel to run the program as though it were being run by the owner of the program.

1. The Set-User-ID Bit

- Doesn't That Mean I Can Change Passwords of Other Users?
 - NO;
 - The `passwd` program knows who you are.
 - It uses the `getuid` system call to ask the kernel for the user ID you used when you logged in.
 - `passwd` has permission to rewrite the entire password file, but will **ONLY** change the **record for the user** running the program.
- Program can test whether a file has SUID bit on by using the mask defined in `<sys/stat.h>`

```
#define S_ISUID          0004000          /* set user id on execution */
```

2. The Set-Group-ID Bit

- The SGID bit sets the effective group ID of a program
 - If a program belongs to group *g* and the set-group-ID bit is set, the program runs as though it were being run by a member of group *g*
- This bit grants the program the access rights of members of that group
- A mask to test for the SGID bit

```
#define S_ISGID          0002000          /* set group id on execution */
```

3. The Sticky Bit

■ Use for files

- In swapping, the sticky bit told the kernel to keep the program on the swap device so that kernel can load it faster. Loading program from swap device is fast because program was never fragmented on the swap device.
- Now, no longer necessary due to virtual memory and paging that allow the kernel to move programs in and out of memory in small sections. We don't need to load entire block of code to run a program

■ Use for directories

- /tmp are publicly writable, allowing any user to create and delete any files there.
- The sticky bit overrides the publicly writable attribute for a directory. Files in the directory may ONLY be deleted by their owners if the sticky bit is set

The Special Bits and ls -l

- Each file has a type and 12 attribute bits, but **ls** uses only 9 spots to display these 12 attributes.

```
-rwsr-sr-t  1 root      root          2345 Jun 12 14:02 sample
```

- Letter **s** indicates that the user and group-executable bits have been augmented by the set-user and set-group ID bits.
- Letter **t** at the end indicates that the sticky bits is on.

Agenda

- What Does **ls** Do?
- Brief Review of the File System Tree
- How **Dos** **ls** Work?
- Can I Write **ls**?
- Writing **ls -l**
- Three Special Bits
- Setting and Modifying the Properties of a File

Type of a File


- A file has a type
 - It can be a regular file(-), a directory(d), a device file(b, c), a socket(s), a symbolic link(l), or a named pipe(p).

- The type of the file is established when the file is created.
 - The creat() system call creates a regular file.
 - Different system call are used to create directories and devices.

- It is not possible to change the type of a file.

Permission Bits and Special Bits

- Every file has 9 permission bits and 3 special bits.



- These bits are established when file is created and can be modified by making the chmod system call

```
fd = creat( "newfile", 0744 );
```

- If you want to prevent programs from creating files that can be modified by group or others
 - `umask(022);`

Permission Bits and Special Bits

\$ man -k chmod

\$ man 2 chmod

■ Changing the mode of a file: chmod() system call

```
chmod( "/tmp/myfile", 04764 );
```

or

```
chmod( "/tmp/myfile", S_ISUID | S_IRWXU | S_IRGRP|S_IWGRP | S_IROTH );
```

■ A Shell Command to Change Permission and Special Bits

```
$ chmod 04764 test
```

or

```
$ chmod u=rws test
```

```
$ chmod g=rw test
```

```
$ chmod o=r test
```

chmod

PURPOSE Change permission and special bits for a file

INCLUDE `#include <sys/types.h>`
 `#include <sys/stat.h>`

USAGE `int result = chmod(char *path, mode_t mode);`

ARGS `path` path to file
 `mode` new value for mode

RETURNS -1 if error
 0 if success

Number of Links to a File

- The number of links is simply the number of times the file is referenced in directories.
 - If a file appears in three places in various directories, the link count is 3. (in the next chapter)

Owner and Group of a File

- Establishing the owner of a file:
 - The owner of file is the user who creates it
 - When kernel creates a file, it sets the owner of the file to be the effective user ID of the process that calls `creat()`
 - If the program has the set-user-ID bit set, though, the effective user ID is the user ID of the person who owns the program.

Owner and Group of a File

- Establishing the group of a file:
 - The group of a file is set to the effective group ID of the process that creates the file.
 - Under non-ordinary circumstances, the group ID of a file is set to the group ID of the parent directory.

Owner and Group of a File

- Changing the owner and group of a File

- chown() system call:

- Normally, users do not change the owner of a file
 - Typically used to set up and manage user accounts

```
chown( "file1", 200, 40 );
```

- Shell Commands to Change User and Group ID for Files:

chown, chgrp

chown

PURPOSE Change owner and or group ID of a file

INCLUDE #include <unistd.h>

USAGE int chown(char *path, uid_t owner, gid_t group)

| | | |
|-------------|-------|-------------------|
| ARGS | path | path to file |
| | owner | user ID for file |
| | group | group ID for file |

| | | |
|----------------|----|------------|
| RETURNS | -1 | if error |
| | 0 | if success |

Modification and Access Time

- Each file has three timestamps of
 - last modified
 - last read
 - file properties (such as owner ID or permission bits) were last changed
 - Kernel automatically updates these times as programs read and write the file

- Changing modification and access times of a file:
 - `utime()` system call

- Shell Commands: `touch`

utime

| | | |
|----------------|---|--|
| PURPOSE | Change access and modification time for files | |
| INCLUDE | #include <sys/time.h> #include <utime.h> | |
| USAGE | #include <sys/types.h> int utime(char *path, struct utimbuf *newtimes) | |
| ARGS | path | path to file |
| | newtimes | pointer to a struct utimbuf see utime.h for details |
| RETURNS | -1 | if error |
| | 0 | if success |

Name of a File

- Establishing the Name of a File

- `creat()` system call sets the name and the initial mode of a file.

- Changing the Name of a File:

- `rename()` system call

- Shell Command : `mv`

- Allows you to change the name of a file
- Also allows you to move a file from one directory to another

rename

PURPOSE Change name and/or move a file

INCLUDE #include <stdio.h>

USAGE int result = rename(char *old, char *new)

ARGS old old name of file or directory
 new new pathname for file or directory

RETURNS -1 if error
 0 if success

VISUAL SUMMARY




FIGURE 3.7

A disk contains files, directories, and their properties.

Objectives

■ Ideas and Skills

- A directory is a list of files
- How to read a directory
- Types of files and how to determine the type of a file
- Properties of files and how to determine properties of a file
- Bit sets and bit masks
- User and group ID numbers

■ System Calls and Functions

- opendir, readdir, closedir, seekdir
- stat
- chmod, chown, utime
- rename

■ Commands

- ls