

Connection Control Studying stty

Prof. Seokin Hong

Kyungpook National University

Fall 2018

Objectives

■ Ideas and Skills

- Similarities between files and devices
- Differences between files and devices
- Attributes of connections
- Race conditions and atomic operations
- Controlling device drivers
- Streams

■ System Calls and Functions

- fcntl, ioctl
- tcsetattr, tcgetattr

■ Commands

- stty
- write

Contents

- 5.2 Devices Are Just Like Files
- 5.3 Devices Are Not Like files
- 5.4 Attributes of Disk Connections
- 5.5 Attributes of Terminal Connections
- 5.6 Programming Other Devices: `ioctl`

Devices Are Just Like Files

- To Unix, sound cards, terminals, mice, and disk files are the same sort of device object.
- In a Unix system, **every device is treated as a file.**
 - Every device has ..
 - a filename
 - an inode #
 - an owner
 - a set of permission bits
 - and a last-modified time
 - Everything working with files automatically applies to terminals and all other devices.
 - We can control devices with system call
 - › open, read, write, lseek, close, stat

Devices Have Filenames

- Files that represent devices are in the **/dev** directory

- Disk file (regular file)

```
add_user.py
[seokin@compasslab1:~$ ls -li test
29361011 -rw-rw-r-- 1 seokin seokin 8  9月 19 14:58 test
seokin@compasslab1:~$
```

inode	mode	link	owner	group	Last_modified time	filename
29361011	-rw-rw-r--	1	seokin	seokin	8 9月 19 14:58	test

- Device file

```
[seokin@compasslab1:~$ ls -il /dev/sda
341 brw-rw---- 1 root disk 8, 0  9月 28 06:16 /dev/sda
seokin@compasslab1:~$
```

inode	mode	link	owner	group	Last_modified time	filename
341	brw-rw----	1	root	disk	8, 0 9月 28 06:16	/dev/sda

major # : Device Driver

minor # : Device

* Directory(d), Symbolic Link(l), Character Device(c), Block Device(b)

Devices and System Calls

- Devices support all file-related system calls:
open, read, write, lseek, close, stat

- Ex) code to read from a disk

```
int fd = open("/dev/sda", O_RDONLY);  
lseek(fd, 0x0111, SEEK_SET);  
read(fd, buffer, 20);
```

→ read data stored in the hard disk this way while bypassing the filesystem.

- Some devices do not support all file operations
 - /dev/input/mice file does not support the write()
 - Terminals support read() and write(), but they do not support lseek().

Ex: Terminals Are Just Like Files

- A terminal is anything that behaves like the classic keyboard and display unit
 - A telnet or ssh window logged in over the Internet behaves like a terminal.
- Command `tty` prints the file name of your terminal
- We can use any file commands and operations with that file:
`cp`, `>`, `mv`, `ln`, `rm`, `cat`, `ls`

\$ tty

/dev/pts/2

\$ cp /etc/motd /dev/pts/2

Today is Monday, we are running low on disk space. Please delete files.
- your sysadmin

\$ who > /dev/pts/2

bruce	pts/2	Jul 17 23:35	(ice.northpole.org)
bruce	pts/3	Jul 18 02:03	(snow.northpole.org)

Properties of Device Files


- Device files have most of the properties disk files (regular file) have.

```
$ ls -li /dev/pts/2
4 crw--w--w- 1 bruce   tty    136, 2 Jul 18 03:25 /dev/pts/2
```

file type major number, minor number

- Major number and minor number**

- The **inode of a device file** stores, not a file size and storage list, but a **pointer to a device driver** in the kernel.
 - Device driver** : a subroutine in the kernel that gets data into and out of a device
 - Major number** specifies which device driver handles the device
 - Minor number** specifies a particular instance of the device



```
[seoking@compasslab1:~]$ ls -li /dev/sd*
341 brw-rw---- 1 root disk 8,  0  9月 28 06:16 /dev/sda
345 brw-rw---- 1 root disk 8,  1  9月 28 06:16 /dev/sda1
344 brw-rw---- 1 root disk 8, 16  9月 28 06:16 /dev/sdb
346 brw-rw---- 1 root disk 8, 17  9月 28 06:16 /dev/sdb1
347 brw-rw---- 1 root disk 8, 18  9月 28 06:16 /dev/sdb2
348 brw-rw---- 1 root disk 8, 19  9月 28 06:16 /dev/sdb3
349 brw-rw---- 1 root disk 8, 20  9月 28 06:16 /dev/sdb4
350 brw-rw---- 1 root disk 8, 21  9月 28 06:16 /dev/sdb5
```

Properties of Device Files

■ Permission Bits

- What do permission bits mean when the file is really a device?
- In the following example,
 - Writing permission means that the owner of the device file and members of group tty are allowed to write to the terminal
 - Reading permission means only owner of the device file can read any terminal messages

```
$ ls -li /dev/pts/2  
4 crw--w--w- 1 bruce tty 136, 2 Jul 18 03:25 /dev/pts/2
```

Writing write

```
$ write test /dev/pts/2
```

- \$ man 1 write

WRITE(1) Linux Programmer's Manual WRITE(1)

NAME

write - send a message to another user

SYNOPSIS

write user [ttyname]

DESCRIPTION

Write allows you to communicate with other users by copying lines from your terminal to theirs.

When you run the write command, the user you are writing to gets a message of the form:

```
Message from yourname@yourhost on yourtty at hh:mm
...
```

```

/* write0.c
 *
 *      purpose: send messages to another terminal
 *      method: open the other terminal for output then
 *              copy from stdin to that terminal
 *      shows: a terminal is just a file supporting regular i/o
 *      usage: write0 ttyname
 */

```

```

#include      <stdio.h>
#include      <fcntl.h>
#include      <stdlib.h>
#include      <string.h>

main( int ac, char *av[] )
{
    int      fd;
    char      buf[BUFSIZ];

    /* check args */
    if ( ac != 2 ){
        fprintf(stderr, "usage: write0 ttyname\n");
        exit(1);
    }

    /* open devices */
    fd = open( av[1], O_WRONLY );
    if ( fd == -1 ){
        perror(av[1]); exit(1);
    }

    /* loop until EOF on input */
    while( fgets(buf, BUFSIZ, stdin) != NULL )
        if ( write(fd, buf, strlen(buf)) == -1 )
            break;

    close( fd );
}

```

```

$ cc write0.c -o write0
$ ./write0 /dev/pts/2
Is it working?...
Bye
^C

```

fgets - get a string from a stream

SYNOPSIS

#include <stdio.h>


**char *fgets(char *restrict s, int n, FILE
*restrict stream);**

DESCRIPTION

The *fgets()* function shall read bytes from *stream* into the array pointed to by *s*, until *n*-1 bytes are read, or a <newline> is read and transferred to *s*, or an end-of-file condition is encountered. The string is then terminated with a null byte.

Device Files and Inodes

- The distinction between types of files appears at the inode level, not in a directory
- **inode** can be a disk file inode or a device-file inode
 - A disk-file inode contains a pointer to the data blocks
 - A device-file inode contains a pointer into a table of kernel subroutines called device drivers



Device Files and Inodes

■ How **read()** works?

- Kernel finds the inode for the file descriptor.
 - inode tells the kernel the type of the file
 - How?
 - › type of inode is recorded into the type portion of the `st_mode` member of `struct stat`
- If the file is a disk file,
 - the kernel gets data by consulting the block allocation list
- If the file is a device file,
 - the kernel reads data by calling the read part of the device driver code for that device.

Contents

- 5.2 Devices Are Just Like Files
- **5.3 Devices Are Not Like files**
- 5.4 Attributes of Disk Connections
- 5.5 Attributes of Terminal Connections
- 5.6 Programming Other Devices: `ioctl`
- 5.7 Up in the Sky!, It's a Device, It's a STREAM!

Devices Are Not Like Files

- Connection to a disk file is different from a connection to a terminal
 - Connection to a disk file usually involves kernel buffers. Buffering is an attribute of the connection
 - Connection to terminal send data to terminal pretty quickly
- We can control attributes in a different way.





FIGURE 5.2

A process with two file descriptors.

Controlling the connection attributes

- You modify the device attributes by changing control variables of **the file descriptor** associated to the device.
- Three steps
 - Get setting
 - Modify them
 - Send them back



To change driver settings:
1. Get settings,
2. modify them
3. send them back.

FIGURE 5.4

Modifying the operation of a file descriptor.

Controlling the connection attributes

■ System call: fcntl

fcntl		
PURPOSE	Control file descriptors	
INCLU D E	#include <fcntl.h> #include <unistd.h> #include <sys/types.h>	
USAGE	int result = fcntl(int fd, int cmd); int result = fcntl(int fd, int cmd, long arg); int result = fcntl(int fd, int cmd, struct flock *lockp);	
ARGS	fd	the file descriptor to control
	cmd	the operation to perform
	arg	arguments to the operation
	lock	lock information
RETURN S	-1	if error
	Other	depends on operation

Controlling the attributes of disk connection

■ Ex1: Turning off disk buffering

```
#include <fcntl.h>
int s;                                // settings
s = fcntl(fd, F_GETFL);               // get flags
s |= O_SYNC;                          // set SYNC bit
result = fcntl(fd, F_SETFL, s);       // set flags
if ( result == -1 )                   // if error
    perror("setting SYNC");           // report
```

- O_SYNC : tells the kernel that calls to write() should be return only when the bytes are written to the actual hardware

Controlling the attributes of disk connection

■ Ex2: Turning on auto-append mode

- With auto-append mode, each call to write() automatically includes an lseek to the end of file

```
#include <fcntl.h>

int s;                                // settings
s = fcntl(fd, F_GETFL);               // get flags
s |= O_APPEND;                        // set APPEND bit
result = fcntl(fd, F_SETFL, s);       // set flags

if ( result == -1 )                   // if error
    perror("setting APPEND");         //      report
else
    write(fd, &rec, 1);               // write record at end
```

Controlling the attributes with open

- `open()` lets you specify fd attribute bits as part of its second argument

```
fd = open(WTMP_FILE, O_WRONLY|O_APPEND|O_SYNC);
```

Summary of disk connections

- The kernel transfers data between disks and processes
- The code in the kernel that performs these transfers has many options
- A program can use the `open()` and `fcntl()` system calls to control the attributes of these data transfers




FIGURE 5.7

Connections to files have settings.

Terminal connection

■ Terminal Driver

- The collection of kernel subroutines that process data flowing between a process and a terminal (external device)
- The driver contains many settings that control its operation
- A process may read, modify, and reset those driver control flags.

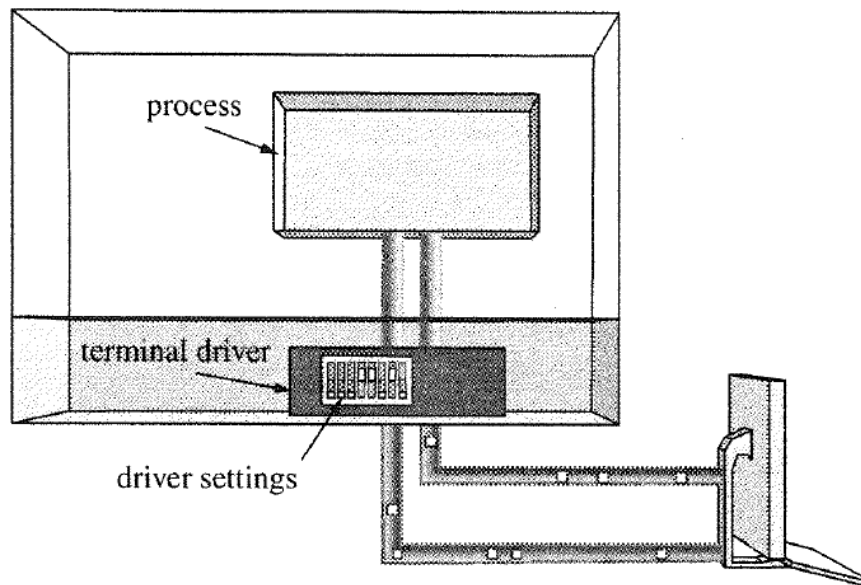


FIGURE 5.10

The terminal driver is part of the kernel.

Controlling the attributes of terminal connection

■ Using stty to Display Driver Settings

```
[seokin2@compasslab1:~$ stty
speed 9600 baud; line = 0;
eol = M-^?; eol2 = M-^?;
-brkint ixany iutf8
-echok
```

■ Using stty to Change Driver Settings

```
$ stty erase X          # make 'X' the erase key
$ stty -echo            # type invisibly  ※ echo off
$ stty erase @ echo   # multiple requests
```

※ echo on

Controlling the attributes of terminal connection

■ Changing setting in the terminal driver

- (a) Get the attributes from the driver.
- (b) Modify any attributes you need to change.
- (c) Send those revised attributes back to the driver.

■ Ex1: Turning on keystroke echoing

```
#include <termios.h>
struct termios settings          /* struct to hold attributes */
tcgetattr(fd, &settings);       /* get attris from driver */
settings.c_lflag |= ECHO;       /* turn on ECHO bit in flagset */
tcsetattr(fd, TCSANOW, &settings); /* send attris back to driver */
```

Ex1: Turning on keystroke echoing

- tcgetattr() library function

tcgetattr

PURPOSE Read attributes from tty driver

INCLUDE `#include <termios.h>`
 `#include <unistd.h>`

USAGE `int result = tcgetattr(int fd, struct termios *info);`

ARGS `fd` file descriptor connected to a terminal
 `info` pointer to a struct termios

RETURNS `-1` if error
 `0` if success

Ex1: Turning on keystroke echoing

■ tcsetattr() library function

tcsetattr							
PURPOSE	Set attributes in tty driver						
INCLUDE	<code>#include <termios.h></code> <code>#include <unistd.h></code>						
USAGE	<code>int result = tcsetattr(int fd, int when, struct termios *info);</code>						
ARGS	<table><tr><td>fd</td><td>file descriptor connected to a terminal</td></tr><tr><td>when</td><td>when to change the settings</td></tr><tr><td>info</td><td>pointer to a struct termios</td></tr></table>	fd	file descriptor connected to a terminal	when	when to change the settings	info	pointer to a struct termios
fd	file descriptor connected to a terminal						
when	when to change the settings						
info	pointer to a struct termios						
RETURNS	<table><tr><td>-1</td><td>if error</td></tr><tr><td>0</td><td>if success</td></tr></table>	-1	if error	0	if success		
-1	if error						
0	if success						

- **TCSANOW** : update driver setting immediately
- **TCSADRAIN** : wait until all output already queued in the driver has been transmitted to the terminal
- **TCSAFLUSH** : wait until all output already queued in the driver has been transmitted. Next, flush all queued input data

Ex1: Turning on keystroke echoing

- The **struct termios** data type contains several flagsets and an array of control characters

```
struct termios
{
    tcflag_t c_iflag;          /* input mode flags */
    tcflag_t c_oflag;          /* output mode flags */
    tcflag_t c_cflag;          /* control mode flags */
    tcflag_t c_lflag;          /* local mode flags */
    cc_t      c_cc[NCCS];      /* control characters */
    speed_t   c_ispeed;        /* input speed */
    speed_t   c_ospeed;        /* output speed */
};
```

Ex1: Turning on keystroke echoing

- Bits and chars in **struct termios** members




FIGURE 5.12
Bits and chars in termios members.

Ex2: Reporting the current state of the terminal

```
/* echostate.c
 *   reports current state of echo bit in tty driver for fd 0
 *   shows how to read attributes from driver and test a bit
 */

#include      <stdio.h>
#include      <termios.h>
#include      <stdlib.h>
main()
{
    struct termios info;
    int rv;
    //'fd = 0' indicates standard input which will be keyboard in the terminal
    rv = tcgetattr( 0, &info );      /* read values from driver      */

    if ( rv == -1 ){
        perror( "tcgetattr" );
        exit(1);
    }
    if ( info.c_lflag & ECHO )      //Test a bit
        printf(" echo is on , since its bit is 1\n");
    else
        printf(" echo if OFF, since its bit is 0\n");
}
```

Ex2: Reporting the current state of the terminal

- Compile and test echostate.c

```
$ cc echostate.c -o echostate
$ ./echostate
echo is on , since its bit is 1
$ stty -echo
$ ./echostate
$ echo is OFF, since its bit is 0
```

Ex2: Turning on/off echo mode

```
/* setecho.c
 *  usage:  setecho [y|n]
 *  shows:  how to read, change, reset tty attributes
 */

#include      <stdio.h>
#include      <termios.h>
#include      <stdlib.h>
#define  oops(s,x) { perror(s); exit(x); }

main(int ac, char *av[])
{
    struct termios info;

    if ( ac == 1 )
        exit(0);
    //fd = 0' indicates standard input which will be keyboard in the terminal
    if ( tcgetattr(0,&info) == -1 )          /* get attribs   */
        oops("tcgetattr", 1);

    if ( av[1][0] == 'y' )
        info.c_lflag |= ECHO ;              /* turn on bit   */
    else
        info.c_lflag &= ~ECHO ;              /* turn off bit  */
        //Set a bit
        //Clear a bit
    if ( tcsetattr(0,TCSANOW,&info) == -1 ) /* set attribs   */
        oops("tcsetattr",2);
}
```


Ex2: Turning on/off echo mode

- Compile and test setecho.c

```
$ echostate; setecho n ; echostate ; stty echo
```

```
echo is on, since its bit is 1
```

```
echo is OFF, since its bit is 0
```

```
$ stty -echo ; echostate ; setecho y ; setecho n
```

```
echo is OFF, since its bit is 0
```

Contents

- 5.2 Devices Are Just Like Files
- 5.3 Devices Are Not Like files
- 5.4 Attributes of Disk Connections
- 5.5 Attributes of Terminal Connections
- 5.6 Programming Other Devices: `ioctl`

Programming Other Devices: ioctl

- What about connections to other types of devices?
 - Ex) CD recorders, Scanners,...

- Every device file supports the ioctl() system call
 - The ioctl() provides access to the attributes and operations of the device driver to fd.
 - Each type of device has its own set of properties and ioctl operations.
 - Ex) A video terminal screen has a size measured in rows and columns or in pixels

```
#include <sys/ioctl.h>

void print_screen_dimensions()
{
    struct winsize wbuf;

    if ( ioctl(0, TIOCGWINSZ, &wbuf) != -1 ){
        printf("%d rows x %d cols\n", wbuf.ws_row, wbuf.ws_col);
        printf("%d wide x %d tall\n", wbuf.ws_xpixel, wbuf.ws_ypixel);
    }
}
```

TIOCGWINGZ is the function code

Programming Other Devices: ioctl

ioctl		
PURPOSE	Control a Device	
INCLUDE	# include <sys/ioctl.h>	
USAGE	int result = ioctl (int fd, int operation [, arg..])	
ARGS	fd	file descriptor connected to device
	operation	operation to perform
	arg...	any args required for the operation
RETURNS	-1	if error
	other	depends on device

VISUAL SUMMARY




FIGURE 5.13

File descriptors, connections, and drivers.